


Center Circle

Project Overview & Guidelines



Change Log

Date	Updates

1. Project Introduction

Center Circle is a product-experience evaluation comparing how leading AI platforms handle the same task over a full conversation. Each task gives you a **preset prompt**, a set of **Considerations for Next Steps**, and a **final goal to reach**. You run that prompt against **two** assigned models, carry each conversation as far as it needs to go (**up to 10 turns**) to reach a fully satisfying answer, score each model on 11 dimensions, write feedback for each, and then pick which model you preferred.

Your scores and written feedback are the data this project depends on.

 **WARNING - NO LLM-WRITTEN FEEDBACK:** Large Language Models are strictly prohibited for writing your feedback or justifications. All written evaluations must be your original analysis. Using AI to generate these responses compromises data integrity and will result in immediate removal from the project.

 **LANGUAGE NOTE:** The prompt is delivered ready-to-paste in your selected language - you do **not** translate it yourself. However, **all feedback and justifications must be written in English.**

2. The Three Task Elements

Every task shows you three things up front. Read all three **before** you run anything - together they tell you what a successful conversation looks like and when to stop.

Element	What it is	How you use it
Prompt	The exact user message, ready to paste in your language.	Copy it into the tool to generate the comparison code, then run it against each model. Do not rewrite it.
Considerations for Next Steps	A short list of plausible directions the user might take the conversation.	Your first follow-up turn must come from this list. After that, you may follow up however best reaches the goal.
Final goal to reach	A plain statement of what a <i>fully satisfying answer</i> means for this task.	This is your stop condition. Keep going until the conversation satisfies it; then stop.

3. Setup - Install the Outlier Playground Extension (*one-time*)

This Chrome extension is your primary interface for session recording and response comparison. You cannot access tasks until installation is verified.

- **Install:** Add the Outlier Playground extension from the Chrome Web Store.
 - **Pin it:** Click the puzzle-piece icon in your Chrome toolbar and pin Outlier Playground for easy access.
 - **Verify:** Confirm the icon is visible in your toolbar before moving on.
-

4. Per-Task Loop

Repeat this for every task you claim.

Setup (steps 1 - 3)

1. **Claim a task** from the queue and open the Outlier Playground extension.
2. **Copy the preset prompt from the task into the tool in outlier**, then run the tool to **generate the comparison code**.
3. **Enter the comparison code** into the "Enter your comparison code" field, click **Continue**, and accept session recording permissions.

Evaluate Model A (steps 4 - 7)

4. **Open Model A** via the extension link in a new tab. Your prompt should auto-fill - if it doesn't, paste it manually. Generate the response.
5. **Run the conversation to a satisfying answer.** Your **first follow-up must pick one of the Considerations for Next Steps**. Continue (up to 10 turns total) until the response satisfies the **final goal to reach** - see *Running the Conversation* below.
6. **Score all 11 dimensions (1–3)** for Model A on the whole conversation, and **write the Model A feedback** (4 - 6 sentences, judged on its own merits - **no comparison to Model B yet**).
7. **Publish the conversation** inside the model platform to generate a shareable link. Return to the Outlier task, paste the link and your feedback into the Model A fields, record which platform Model A was and continue in the outlier extension.

Evaluate Model B (steps 8–11)

Repeat the Model A sequence with the second assigned model, **using the same starting prompt**:

8. **Open Model B** via the extension link. Generate the response.
9. **Run the conversation** to a satisfying answer, applying the **same patience** you gave Model A (see parity rule below).
10. **Score all 11 dimensions (1–3)** and **write the Model B feedback** independently - judge it on its own merits, not against Model A yet.

11. **Publish the conversation**, paste the link and feedback into the Model B fields, record which platform Model B was, and click **"Done with conversation."**

Compare and submit (steps 12–14)

12. **Paste in selected next step:** When asked in the task, copy and paste in the next step that you selected to take the model in the direction of from the considerations for next steps.
13. **Pick your preference:** *Model A better, Model B better, I like both, or neither.* There is no "tie" — "like both" and "dislike both" are different signals. Then **write your overall justification** (4–6 sentences) 0 this is the **only** place comparison belongs.
14. **Submit** the comparison in the extension and submit your task after ensuring you pasted in all relevant information!

5. Running the Conversation - Turns & Stop Condition

Tasks are open-ended on length. Go as many turns as you need (max **10**) to reach a fully satisfying answer according to the **final goal to reach**, then stop. If the first response already nails it, you may still need at least your one required follow-up to test how the model handles a continuation - see below.

Three rules:

1. **Your first follow-up must come from the Considerations for Next Steps.** Pick the one that best moves the conversation toward the goal. You do not need to use all of them - one is required.
2. **After the first follow-up, you're free.** If the model takes the conversation in a different direction, or a different consideration (or a brand-new follow-up you write yourself) gets you to the goal faster, that's completely fine. Follow-ups should **respond to what the model actually said** - engage with the real response rather than re-pasting or rephrasing the original prompt.
3. **Treat both models the same (parity).** Whatever patience you extend to Model A, extend to Model B. Pushing one model through several follow-ups while abandoning the other after a weak first turn creates unfair signal. How quickly each model reached a satisfying answer is real signal - capture it in your feedback and let it inform your scores.

6. Dimensions Scoring

For **each model**, score **all 11 dimensions** on a **1–3 scale**, judging the **whole conversation** (not any single turn).

Score Model A and Model B **independently**. Do not let one model's score influence the other's - the comparison happens only at the final preference step.

The 11 dimensions fall into two groups: **Model Quality** (1–7) and **Product Experience** (8–11).

Model Quality Rubrics

Dimension	3 — No Issue	2 — Minor Issue	1 — Major Issue
1. Task Completion <i>Did the model help the user accomplish their goal? Whether the response meaningfully addresses the request, resolves the main intent, and moves toward completion.</i>	Fully satisfies the user's objective with little or no additional effort required.	Mostly satisfies the objective, but has small gaps, inefficiencies, or missing details that don't prevent the user from moving forward.	Only partially addresses the objective, misunderstands the core request, or leaves the user needing significant additional work.
2. Instruction Following <i>Did the model follow the user's explicit and implicit instructions? Includes format, tone, constraints, language, locale, length, style, and any requested do's or don'ts.</i>	Follows all explicit and implicit instructions, including format, tone, constraints, and subtle preferences.	Follows the main instructions, but misses one or more smaller constraints that don't significantly affect usefulness.	Ignores, contradicts, or fails to follow important instructions, making the output hard to use or misaligned with the request.
3. Accuracy and Groundedness <i>Is the response factually correct, reliable, and appropriately grounded? Covers factual accuracy, logical correctness, use of available context, and avoidance of unsupported claims.</i>	Accurate, reliable, well-supported, and appropriately caveated where needed.	Mostly accurate, but includes small unsupported claims, minor imprecision, or caveats that could be clearer.	Contains clear factual errors, hallucinations, misleading claims, or unsupported assertions that materially affect quality.
4. Relevance and Focus <i>Is the response focused on what the user actually asked? Whether the</i>	Directly focused on the request; avoids unnecessary or irrelevant content.	Mostly relevant, but includes some extra, generic, or	Largely off-topic, misses the actual intent, or focuses on

Dimension	3 — No Issue	2 — Minor Issue	1 — Major Issue
model avoids unnecessary tangents, over-answering, boilerplate, or generic content.		slightly unfocused content.	the wrong parts of the request.
5. Clarity and Communication Quality <i>Is the response easy to understand and well structured?</i> Organization, readability, concision, wording, and whether it feels human-written rather than mechanical.	Clear, natural, well organized, concise, and easy to act on.	Understandable, but may be slightly wordy, awkward, poorly organized, or harder to scan than needed.	Confusing, poorly structured, overly verbose, unnatural, or difficult to understand or use.
6. Helpfulness and Practicality <i>Is the response genuinely useful in the user's real-world context?</i> Whether it gives actionable, practical, context-aware help rather than generic advice.	Highly useful, actionable, and tailored to the user's real-world context.	Useful overall, but could be more actionable, specific, or tailored to the situation.	Generic, impractical, hard to apply, or not meaningfully useful.
7. Locale and Cultural Fit <i>Does the response fit the user's language, locale, and cultural context?</i> Covers correct language/dialect, local terminology, cultural norms, units/currency/dates/form atting, region-specific expectations, and avoidance of unnatural translations. Especially important for multilingual and region-specific evals.	Feels native, locally appropriate, and aligned with the user's language, region, and cultural context.	Understandable and mostly localized, but with small issues in phrasing, terminology, formatting, or cultural fit.	Uses the wrong language or dialect, feels unnatural for the locale, applies incorrect regional assumptions, or is culturally inappropriate.

Product Experience Rubrics

Dimension	3 — No Issue	2 — Minor Issue	1 — Major Issue
8. Conversation Flow <i>Does the assistant maintain a smooth, useful conversation over multiple turns?</i> Whether it remembers context, builds on prior turns, avoids making the user repeat themselves, and advances naturally.	Maintains context smoothly, builds on prior turns, and moves the interaction forward naturally.	Mostly maintains flow, but has small moments of repetition, missed nuance, or slightly awkward turn-taking.	Loses track of the conversation, forces unnecessary repetition, or fails to advance the interaction usefully.
9. Friction and Effort Required <i>How much extra work does the user need to do after receiving the response?</i> Whether the experience feels efficient or frustrating.	The user can act immediately with little or no additional clarification, editing, or effort.	The user needs small clarifications, edits, or adjustments before the response is usable.	The user needs significant additional work, must re-explain the request, heavily correct the output, or restart the task.
10. Multi-Turn Consistency <i>Does the assistant consistently maintain and apply the user's context, preferences, and intent across the conversation?</i> Rewards using relevant info from earlier turns, continuity in tone and direction, and not making the user repeat themselves. Penalizes contradicting prior context, forgetting stated preferences, or over-assuming.	Consistently carries forward context, preferences, constraints, and goals across turns in a helpful, natural way.	Maintains the main context and intent, but misses some smaller prior details or shows minor gaps in continuity.	Forgets or contradicts important prior context, ignores stated preferences/constraints, makes inappropriate assumptions, or requires the user to repeat key information.
11. Delight and Overall Experience	Polished, natural, satisfying, and pleasant;	Usable and generally pleasant, but may feel flat,	Frustrating, awkward, unpleasant, or

Dimension	3 — No Issue	2 — Minor Issue	1 — Major Issue
<i>Does the interaction feel satisfying, polished, and pleasant?</i> The softer product-quality layer: tone, empathy, naturalness, and whether the assistant feels like something users would want to keep using.	appropriate tone; easy to continue using.	robotic, impersonal, or slightly rough around the edges.	meaningfully below the quality users expect from a helpful assistant.

7. How to Write Strong Feedback

You write **two** feedback blocks (one per model, independent) plus **one** final justification (comparison). **Ensure that in the feedback you write you are talking about how the model behaved in response to the next steps you gave it alongside if the final goal was reached.**

Element	Guidance
Length	4 - 6 sentences. Long enough to dive into what worked and what didn't, short enough to stay focused.
Specificity	Cite concrete examples from the conversation. "The third turn contradicted the constraint set in turn one" beats "consistency was off."
Anchor to the criteria	Tie your points to the dimensions above - your written feedback and your 1–3 scores should tell the same story. If you scored a dimension a 1, the feedback should explain why.
Consideration of Next Steps	Talk about how the model behaved and what worked and what didn't work when you attempted to guide it using one of the next steps provided. Reference specific examples from the model response.
Turn behavior	Note how many turns it took and how the model handled follow-ups - how quickly it reached the final goal is real signal.

Element	Guidance
Independence (Model A & B feedback)	Judge each response on its own merits. Do not write "Model A was better than Model B" in the per-model feedback - save all comparison for the final justification.
Comparison (final justification only)	The final justification is the only place you compare. Name the concrete differences that drove your preference, and acknowledge where the model you didn't pick was actually stronger.
FINAL GOAL	When writing your justifications, you MUST address whether the model reached the final goal, how fully it got there, and what worked or fell short along the way.

Example:

Prompt:

Can you write an artist bio about Kim Myhre Clarke using the information I provide?

Considerations for Next Steps:

1. The user might ask the AI to elaborate on a specific aspect of Kim's artistic process, such as her use of glazes or engobes.
2. The user might request the AI to highlight Kim's world travels and how they influenced her work.
3. The user might ask the AI to focus more on the influence of specific artists like Frida Kahlo or Tim Burton on Kim's style.
4. The user might ask the AI to write a shorter or longer version of the bio for different purposes.
5. The user might ask the AI to suggest ways to incorporate images of Kim's work into the bio.

Final Goal to Reach:

The user wants the AI to write a compelling artist bio for Kim Myhre Clarke, incorporating details about her life, artistic background, influences, creative process, and recent work, based on the provided text.

Example Strong per-model feedback (independent)

The model reached the final goal by the second turn, producing a compelling bio that pulled Kim's background, influences, and creative process straight from the text I provided rather than inventing details (Task Completion, Accuracy and Groundedness). For my first follow-up I used one of the next-step considerations and asked it to highlight her world travels; it wove that in cleanly without dropping the earlier detail about her glaze and engobe work (Multi-Turn Consistency). The structure was clean and the tone read like a professional artist bio, so it was usable with little editing (Clarity, Instruction Following). The one weakness was a generic closing paragraph that added length without new information, slightly diluting the focus (Relevance and Focus).

Why this works: it opens by stating whether and how well the final goal was met, cites specific detail from the actual conversation (the travel follow-up, the glaze work, the weak closing), anchors each point to a named dimension, notes the turn count and the required Considerations follow-up, and stays fully independent - no mention of the other model. Four to five tight sentences, not a paragraph.

Example Strong final justification (comparison)

Model A reached a fully satisfying bio in two turns, while Model B needed four and still left the creative-process section thin, so Model A met the final goal more completely. Model A also kept the bio firmly grounded in the text I provided, whereas Model B invented an exhibition that wasn't in my notes, which undercut its accuracy. Model B did handle the world-travel follow-up with richer, more vivid detail, which was a genuine strength. But the user's goal was a compelling bio built faithfully from the provided information, and Model A delivered that more reliably and with less back-and-forth. For its stronger groundedness and faster goal completion, I preferred Model A.

Why this works: comparison lives only here (never in the per-model feedback), it names the concrete differences that drove the choice, ties the decision back to the user's actual final goal, and honestly credits where the model that wasn't chosen was stronger.

Weak feedback (avoid)

Pretty good response. Helpful and well-written. Better than the other one.

Why it fails: vague with no specifics from the conversation, never states whether or how well the final goal was reached, not anchored to any dimension, and comparative where it should be independent.

8. Common Mistakes to Avoid

- **Vague feedback.** "Good response" tells us nothing. Cite specifics from the conversation.
- **Comparing in the wrong place.** Model A and Model B feedback must be independent. Comparison belongs only in the final justification.
- **Unequal patience.** Running one model through several follow-ups while cutting the other short breaks parity. Treat both the same.
- **Skipping the required first follow-up.** Your first follow-up must come from the Considerations for Next Steps.
- **Stopping before / overshooting the goal.** Use the **final goal** as your stop condition - don't quit early, and don't pad with extra turns once it's satisfied.
- **Scores that don't match the feedback.** If you scored a dimension low, your written feedback should say why. They must agree.
- **Writing feedback in the wrong language.** The prompt is in your language; **all feedback and justifications are in English.**
- **Editing the preset prompt.** Paste it as delivered.